

# **TRIPLEO**

## **PROVISION YOUR DATACENTER WITH OPENSTACK**

Imre Farkas and Ladislav Smola

# AGENDA

- TripleO tools
- How to deploy OpenStack?
- Architecture
- Tuskar & Tuskar-UI
- Advanced features

# WHAT IS TRIPLEO?

---

*“TripleO is a program aimed at installing, upgrading and operating OpenStack clouds using OpenStack's own cloud facilities as the foundations - building on Nova, Neutron, Heat and Ironic to automate fleet management at datacenter scale (and scaling down to as few as 2 machines).”*

# HOW TO LAUNCH AN APPLICATION?

1. grab an image
2. deploy
3. configure

# GRAB AN IMAGE

```
wget http://download.fedoraproject.org/pub/fedora/linux/releases/20/  
      Images/x86_64/Fedora-x86_64-20-20131211.1-sda.qcow2
```

# HOW TO LAUNCH AN APPLICATION?

1. grab an image
2. **deploy**
3. configure

# HOW TO DEPLOY A VM? - SOLUTION #1

```
nova boot myCentOSServer  
  --image "3afe97b2-26dc-49c5-a2cc-a2fc8d80c001"  
  --flavor m1.small
```

Use Puppet, Chef, Ansible or whatever for configuration

# HOW TO DEPLOY A VM? - SOLUTION #2

Use Heat!

# HEAT TEMPLATE

```
heat_template_version: 2013-05-23

description:
  ...

parameters:
  ...

resources:
  ...

outputs:
  ...
```

# HEAT TEMPLATE - PARAMETERS

```
parameters:
  db_name:
    type: string
    description: WordPress database name
    default: wordpress
    constraints:
      - length: { min: 1, max: 64 }
        description: db_name must be between 1 and 64 characters
      - allowed_pattern: '[a-zA-Z][a-zA-Z0-9]*'
        description: >
          db_name must begin with a letter and contain only alphanumeric
          characters
```

# HEAT TEMPLATE - RESOURCES

```
resources:
  wordpress_instance:
    type: OS::Nova::Server
    properties:
      image: { get_param: image_id }
      flavor: { get_param: instance_type }
      key_name: { get_param: key_name }
      user_data:
        ...
...
```

# HEAT TEMPLATE - RESOURCE PROPERTIES: USER\_DATA (1)

```
user_data:
  str_replace:
    template: |
      #!/bin/bash -v

      yum -y install mysql mysql-server httpd wordpress
      systemctl enable mysqld.service
      systemctl enable httpd.service
      systemctl start mysqld.service
      systemctl start httpd.service

      firewall-cmd --add-service=http
      firewall-cmd --permanent --add-service=http
```

# HEAT TEMPLATE - OUTPUTS

```
outputs:
  WebsiteURL:
    description: URL for Wordpress wiki
    value:
      str_replace:
        template: http://host/wordpress
        params:
          host: { get_attr: [wordpress_instance, first_address] }
```

# HOW TO DEPLOY AN VM? - SOLUTION #2

```
heat stack-create mystack  
    --template-file=WordPress_Single_Instance.yaml  
    --parameters="db_name=$db_name;db_rootpassword=..."
```

# **TRIPLEO-HEAT-TEMPLATES**

# HOW TO LAUNCH AN APPLICATION?

1. grab/build an image
2. deploy
3. configure

# **IMAGE BUILDING**

# THE GOAL OF IMAGE BUILDING IN TRIPLEO

---

*“The goal of the image building process is to produce blank slate machines that have all the necessary bits to fulfill a specific purpose in the running of an Openstack cloud: e.g. a nova-compute node.”*

# DISKIMAGE-BUILDER

- disk images / file system images / ramdisk images
- virtual / baremetal machines
- stock of elements
- easily extensible

# DISKIMAGE-BUILDER

## SCRIPTS EXECUTION PHASES

- root.d
- finalise.d
- cleanup.d
- block-device.d
- extra-data.d
- pre-install.d
- install.d
- post-install.d
- environment.d

# ELEMENT EXAMPLE: FEDORA

```
bin/install-packages
bin/map-packages
bin/map-services
finalise.d/01-clean-old-kernels.sh
finalise.d/99-setup-first-boot
install.d/00-fedora-fixup-audit
install.d/00-fedora-fixup-openssl
install.d/00-fedora-fixup-pyopenssl
install.d/01-install-deps
pre-install.d/15-fedora-remove-grub
pre-install.d/00/usr-local-bin-secure-path
pre-install.d/02-lsb
root.d/10-fedora-cloud-image
README.md
element-deps
source-repository-fedora
```

# HOW TO LAUNCH AN APPLICATION?

1. grab/build an image
2. deploy
3. configure

# **CONFIGURATION**

# OS-COLLECT-CONFIG

```
[default]
command=os-refresh-config

[cfn]
metadata_url=http://192.0.2.99:8000/v1/
access_key_id = ABCDEFGHIJLMNOP01234567890
secret_access_key = 01234567890ABCDEFGHIJKLMNP
path = MyResource
stack_name = my.stack
```

# HEAT TEMPLATE - METADATA

```
resources:  
  wordpress_instance:  
    type: OS::Nova::Server  
    properties:  
      image: { get_param: image_id }  
      ...  
    metadata:  
      key: value
```

# OS-REFRESH-CONFIG

## SCRIPTS EXECUTION PHASES

- pre-configure.d
- configure.d
- migration.d
- post-configure.d

# OS-APPLY-CONFIG

converts JSON file to service config

```
{"keystone": {"database": {"host": "127.0.0.1",  
                         "user": "keystone",  
                         "password": "foobar"}}}
```

```
[sql]  
connection = mysql://keystone:foobar@127.0.0.1/keystone
```

# MUSTACHE TEMPLATE EXAMPLE

```
[database]
connection={{tuskar.db}}

[heat_keystone]
username = {{tuskar.user}}
tenant_name = {{tuskar.tenant_name}}
password = {{tuskar.password}}
auth_url = http://{{keystone.host}}:35357/v2.0
```

# TRIPLEO-IMAGE-ELEMENTS

# **DISKIMAGE-BUILDER/ELEMENTS VS TRIPLEO- IMAGE-ELEMENTS**

# ELEMENT EXAMPLE: NOVA

```
install.d/nova-source-install/74-nova
os-apply-config/etc/nova/nova.conf
os-refresh-config/configure.d/10-nova-state
pre-install.d/00-disable-requiretty
README.md
element-deps
source-repository-nova
```

# VM VS BAREMETAL

- pxe boot
- autodiscovery
- hw specs
- perf metrics

# IRONIC

# **TRIPLEO == OPENSTACK ON OPENSTACK**

**IN REALITY: TRIPLEO == OPENSTACK ON  
OPENSTACK ON OPENSTACK**

# ARCHITECTURE

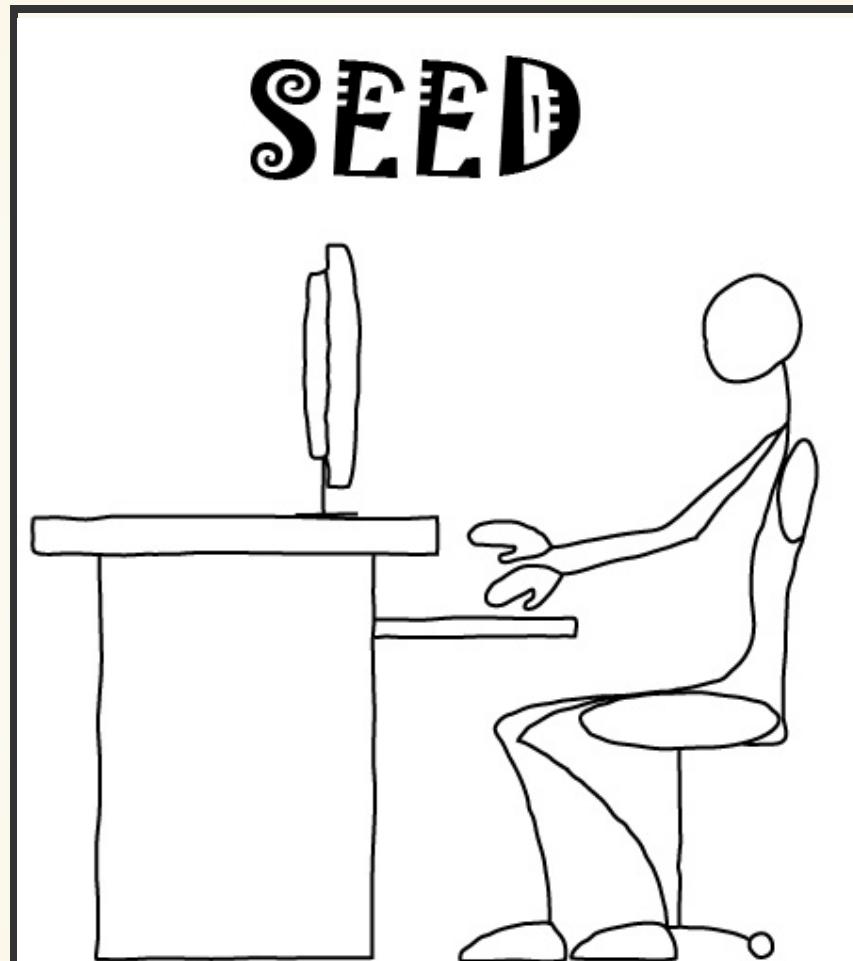
# OVERCLOUD

# UNDERCLOUD

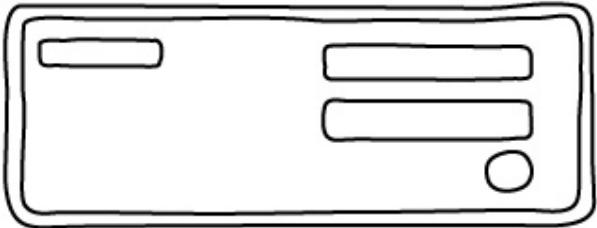
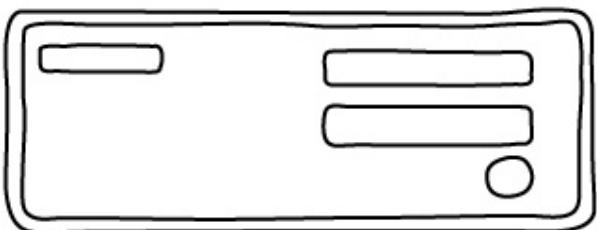
**SEED**

# WORKFLOW

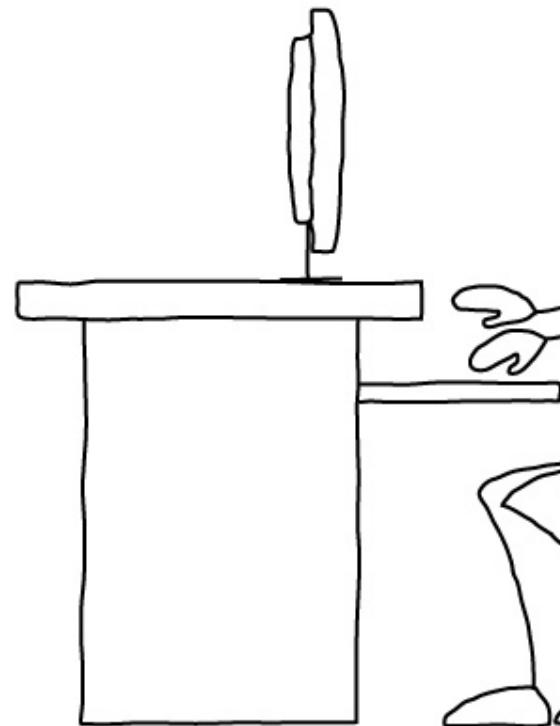
# SEED



# DEPLOY UNDERCLOUD



Undercloud



# DEPLOY OVERCLOUD

# DEVTEST

# SEED - SETUP VM

```
setup-seed-vm -a $NODE_ARCH  
  
$TRIPLEO_ROOT/diskimage-builder/bin/ramdisk-image-create \  
-a $NODE_ARCH $NODE_DIST $DEPLOY_IMAGE_ELEMENT \  
-o $TRIPLEO_ROOT/deploy-ramdisk  
  
boot-seed-vm -a $NODE_ARCH $NODE_DIST neutron-dhcp-agent
```

# SEED - REGISTER SERVICES

```
init-keystone -p unset unset 192.0.2.1 admin@example.com root@192.0.2.1  
setup-endpoints 192.0.2.1 --glance-password unset  
                --heat-password unset  
                --neutron-password unset  
                --nova-password unset
```

# UNDERCLOUD - CREATE IMAGE

```
$TRIPLEO_ROOT/diskimage-builder/bin/disk-image-create $NODE_DIST  
-a $NODE_ARCH -o $TRIPLEO_ROOT/undercloud  
boot-stack nova-baremetal os-collect-config dhcp-all-interfaces  
neutron-dhcp-agent  
  
UNDERCLOUD_ID=$(load-image $TRIPLEO_ROOT/undercloud.qcow2)
```

# UNDERCLOUD - DEPLOY WITH HEAT

```
make -C $TRIPLEO_ROOT/tripleo-heat-templates undercloud-vm.yaml  
  
heat stack-create  
-f $TRIPLEO_ROOT/tripleo-heat-templates/undercloud-vm.yaml  
-P "PowerUserName=$(whoami);AdminToken=${UNDERCLOUD_ADMIN_TOKEN};  
    AdminPassword=${UNDERCLOUD_ADMIN_PASSWORD};  
    GlancePassword=${UNDERCLOUD_GLANCE_PASSWORD};  
    HeatPassword=${UNDERCLOUD_HEAT_PASSWORD};  
    NeutronPassword=${UNDERCLOUD_NEUTRON_PASSWORD};  
    NovaPassword=${UNDERCLOUD_NOVA_PASSWORD};  
    BaremetalArch=${NODE_ARCH};  
    PowerManager=$POWER_MANAGER;  
    undercloudImage=${UNDERCLOUD_ID}"  
undercloud
```

# UNDERCLOUD - REGISTER SERVICES

```
init-keystone -p $UNDERCLOUD_ADMIN_PASSWORD $UNDERCLOUD_ADMIN_TOKEN  
$UNDERCLOUD_IP admin@example.com heat-admin@$UNDERCLOUD_IP  
  
setup-endpoints $UNDERCLOUD_IP  
    --glance-password $UNDERCLOUD_GLANCE_PASSWORD  
    --heat-password $UNDERCLOUD_HEAT_PASSWORD  
    --neutron-password $UNDERCLOUD_NEUTRON_PASSWORD  
    --nova-password $UNDERCLOUD_NOVA_PASSWORD
```

# OVERCLOUD - CREATE THE CONTROL IMAGE

```
$TRIPLEO_ROOT/diskimage-builder/bin/disk-image-create $NODE_DIST  
-a $NODE_ARCH -o $TRIPLEO_ROOT/overcloud-control  
boot-stack cinder-api cinder-volume os-collect-config  
neutron-network-node dhcp-all-interfaces swift-proxy swift-storage  
  
OVERCLOUD_CONTROL_ID=$(load-image -d $TRIPLEO_ROOT/overcloud-control.qcow2)
```

# OVERCLOUD - CREATE THE COMPUTE IMAGE

```
$TRIPLEO_ROOT/diskimage-builder/bin/disk-image-create $NODE_DIST \
-a $NODE_ARCH -o $TRIPLEO_ROOT/overcloud-compute \
nova-compute nova-kvm neutron-openvswitch-agent os-collect-config \
dhcp-all-interfaces

OVERCLOUD_COMPUTE_ID=$(load-image -d $TRIPLEO_ROOT/overcloud-compute.qcow2)
```

# UNDERCLOUD - DEPLOY WITH HEAT

```
make -C $TRIPLEO_ROOT/tripleo-heat-templates overcloud.yaml  
  
heat stack-create -f $TRIPLEO_ROOT/tripleo-heat-templates/overcloud.yaml  
-P "AdminToken=${OVERCLOUD_ADMIN_TOKEN};  
    AdminPassword=${OVERCLOUD_ADMIN_PASSWORD};  
    CinderPassword=${OVERCLOUD_CINDER_PASSWORD};  
    GlancePassword=${OVERCLOUD_GLANCE_PASSWORD};  
    HeatPassword=${OVERCLOUD_HEAT_PASSWORD};  
    NeutronPassword=${OVERCLOUD_NEUTRON_PASSWORD};  
    NovaPassword=${OVERCLOUD_NOVA_PASSWORD};  
    NeutronPublicInterface=${NeutronPublicInterface};  
    SwiftPassword=${OVERCLOUD_SWIFT_PASSWORD};  
    SwiftHashSuffix=${OVERCLOUD_SWIFT_HASH}${OVERCLOUD_LIBVIRT_TYPE};  
    SSLCertificate=${OVERCLOUD_SSL_CERT};SSLKey=${OVERCLOUD_SSL_KEY};  
overcloud
```

# UNDERCLOUD - REGISTER SERVICES

```
init-keystone -p $OVERCLOUD_ADMIN_PASSWORD $OVERCLOUD_ADMIN_TOKEN  
$OVERCLOUD_IP admin@example.com heat-admin@$OVERCLOUD_IP  
${SSLBASE:+--ssl $PUBLIC_API_URL}

setup-endpoints $OVERCLOUD_IP --cinder-password $OVERCLOUD_CINDER_PASSWORD  
--glance-password $OVERCLOUD_GLANCE_PASSWORD  
--heat-password $OVERCLOUD_HEAT_PASSWORD  
--neutron-password $OVERCLOUD_NEUTRON_PASSWORD  
--nova-password $OVERCLOUD_NOVA_PASSWORD  
--swift-password $OVERCLOUD_SWIFT_PASSWORD  
${SSLBASE:+--ssl $PUBLIC_API_URL}
```

# CI/CD PIPELINE

# UPDATES

# SO FAR...

- TripleO tools
- How to deploy OpenStack?
- Architecture

# **OVERCLOUD DEPLOYMENT WITH TUSKAR**

# REGISTERING OF BAREMETAL NODES(MANUAL)

**Register Nodes**

5 Nodes to Register    [+ Add Node](#)

192.168.0.1
192.168.0.2
192.168.0.3
192.168.0.4
192.168.0.5

**Node Detail**

Node Tags

power socket P1 ×

my group × hp-comp ×

Management

IP Address \*

IPMI User

IPMI Password

Networking

NIC MAC Address \*

01:23:45:67:89:a1

01:23:45:67:89:a2

Hardware

Architecture \*

CPUs \*

cores

Memory \*

GB

Local Disk \*

TB

# HARDWARE PROFILES (FLAVORS)

- Standard openstack flavors with architecture field (i386/amd64) added
- Heat uses nova-boot for provisioning, nova scheduler looks for available hardware by doing exact match of node specs with chosen flavor

# IMAGES

- Created by image builder and uploaded to Glance.

# DEPLOYMENT ROLE

- Creates a group of baremetals running the same group of services. (E.g. nova compute, block storage, etc.)
- Relation to hardware profiles(flavors)
- Relation to image
- Relation to heat template

# LIMITED FOR THESE ROLES FOR ICEHOUSE

- Controller (all services control planes)
- Compute (Nova)
- Object storage (Swift)
- Block storage (Cinder)

# DEPLOYMENT ROLE EXAMPLE

## Edit Deployment Role

General Settings

Name      Controller

Provisioning Image

Image      controller-image

Node Profiles

- My Profile | x86\_64 | 2 CPU cores | 4 GB RAM | 10 TB HDD | No Tags
- Another Node Profile | i386 | 1 CPU core | 6 GB RAM | 1 TB HDD | No Tags
- Different Profile | i386 | 1 CPU core | 6 GB RAM | 1 TB HDD | No Tags
-

# DEPLOYING OVERCLOUD EXAMPLE

Deployment Roles		65 free nodes	Configuration
Name	Node profiles	Instances	
 Controller	Profile 1	<div><span>1</span></div>	
 Compute	My Node Profile	<div><span>50</span></div>	
	Different Profile	<div><span>10</span></div>	
 Object Storage	Profile 2	<div><span>4</span></div>	
 Block Storage	Profile 2	<div><span>0</span></div>	

# DEPLOYING OVERCLOUD EXAMPLE

## My OpenStack Deployment

[Delete](#) [Scale Deployment](#)

[Overview](#) [Configuration](#) [Log](#)

Deploying...  Last update:  
02:03:04, 2014-01-16 UTC: overcloud\_compute0\_osij2947asj2 is being created. [See full log](#)

---

### Deployment Roles

Controller	 0 / 1 instances has been deployed
Compute	 1 / 50 instances has been deployed
Object Storage	 1 / 4 instances has been deployed
Block Storage	-

# DEPLOYING OVERCLOUD EXAMPLE

## My OpenStack Deployment

Overview Configuration Log Undeploy Scale Deploy

Deployment Roles

<a href="#">Controller (1)</a>	<span>✓ All instances run correctly</span>	 93 %
<a href="#">Compute (60)</a>	<span>⚠ 1 instance is down</span>	 65 %
<a href="#">Object Storage (4)</a>	<span>✓ All instances run correctly</span>	 21 %
<a href="#">Block Storage (0)</a>	<span>- No instance</span>	

Deployment Roles Distribution



Role	Percentage
Controller	8 %
Object Storage	25 %
Compute	67 %

Horizon UI Connection

[Horizon UI \(<http://192.168.11.5>\)](#)

# FUTURE (JUNO) MANAGEMENT OF DEPLOYMENT ROLES

- Assigning Image
- Assigning Template
- Assigning Hardware profiles (flavors)

Then we are prepared to deploy any service on the separate baremetals, that can be easily scaled, e.g. Neutron, Ceilometer, Ironic, etc....

In future we will allow to provision multiple images on one machine with Docker.

# FUTURE POSSIBLE NEEDED SERVICES

- Heat Template repository as Openstack service?
- Image builder as a Openstack service?

# OVERCLOUD MONITORING WITH TUSKAR

# MONITORING OF BAREMETALS

- Icehouse: monitoring via SNMP: cpu\_util, memory, disk, network
- Juno: monitoring via IPMI: Temperature, Fan Speed, Volt, etc. Plus pluggable architecture for vendor specific metrics.

# STORING AND QUERYING STATS

- Using Ceilometer agent for polling of the samples
- Using Ceilometer as generic samples and meters storage
- Using Ceilometer for querying statistics
- Ceilometer is supporting many backends MongoDB, MySQL, PostgreSQL, HBase, DB2 (Elastic search in progress)

# USING D3 AND RICKSHAW LIBRARIES FOR RENDERING CHARTS

Deployment Role: Controller

5 instances

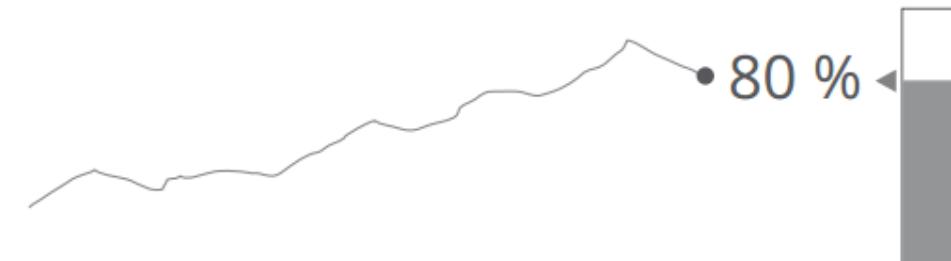
2 node profiles

5x My node profile: 1 CPU | 16 GB RAM | 4 TB HDD  
0x Other profile: 2 CPU | 8 GB RAM | 1 TB HDD

controller-image

5 Nodes

	Node	Node Profile	Capacity	Architecture	CPUs (cores)	Memory (GB)	Disk (TB)	Health	Power	Actions
<input type="checkbox"/>	192.168.0.1	My node profile	<div style="width: 100px; height: 10px; background-color: #ccc;"></div>	x86_64	1	16	4	Operational	On	<a href="#">Actions</a>
<input type="checkbox"/>	192.168.0.2	My node profile	<div style="width: 100px; height: 10px; background-color: #ccc;"></div>	x86_64	1	16	4	Operational	On	<a href="#">Actions</a>
<input type="checkbox"/>	192.168.0.3	My node profile	<div style="width: 100px; height: 10px; background-color: #ccc;"></div>	x86_64	1	16	4	Operational	On	<a href="#">Actions</a>
<input type="checkbox"/>	192.168.0.4	My node profile	<div style="width: 100px; height: 10px; background-color: #ccc;"></div>	x86_64	1	16	4	Operational	On	<a href="#">Actions</a>
<input type="checkbox"/>	192.168.0.5	My node profile	<div style="width: 100px; height: 10px; background-color: #ccc;"></div>	x86_64	1	16	4	Operational	On	<a href="#">Actions</a>



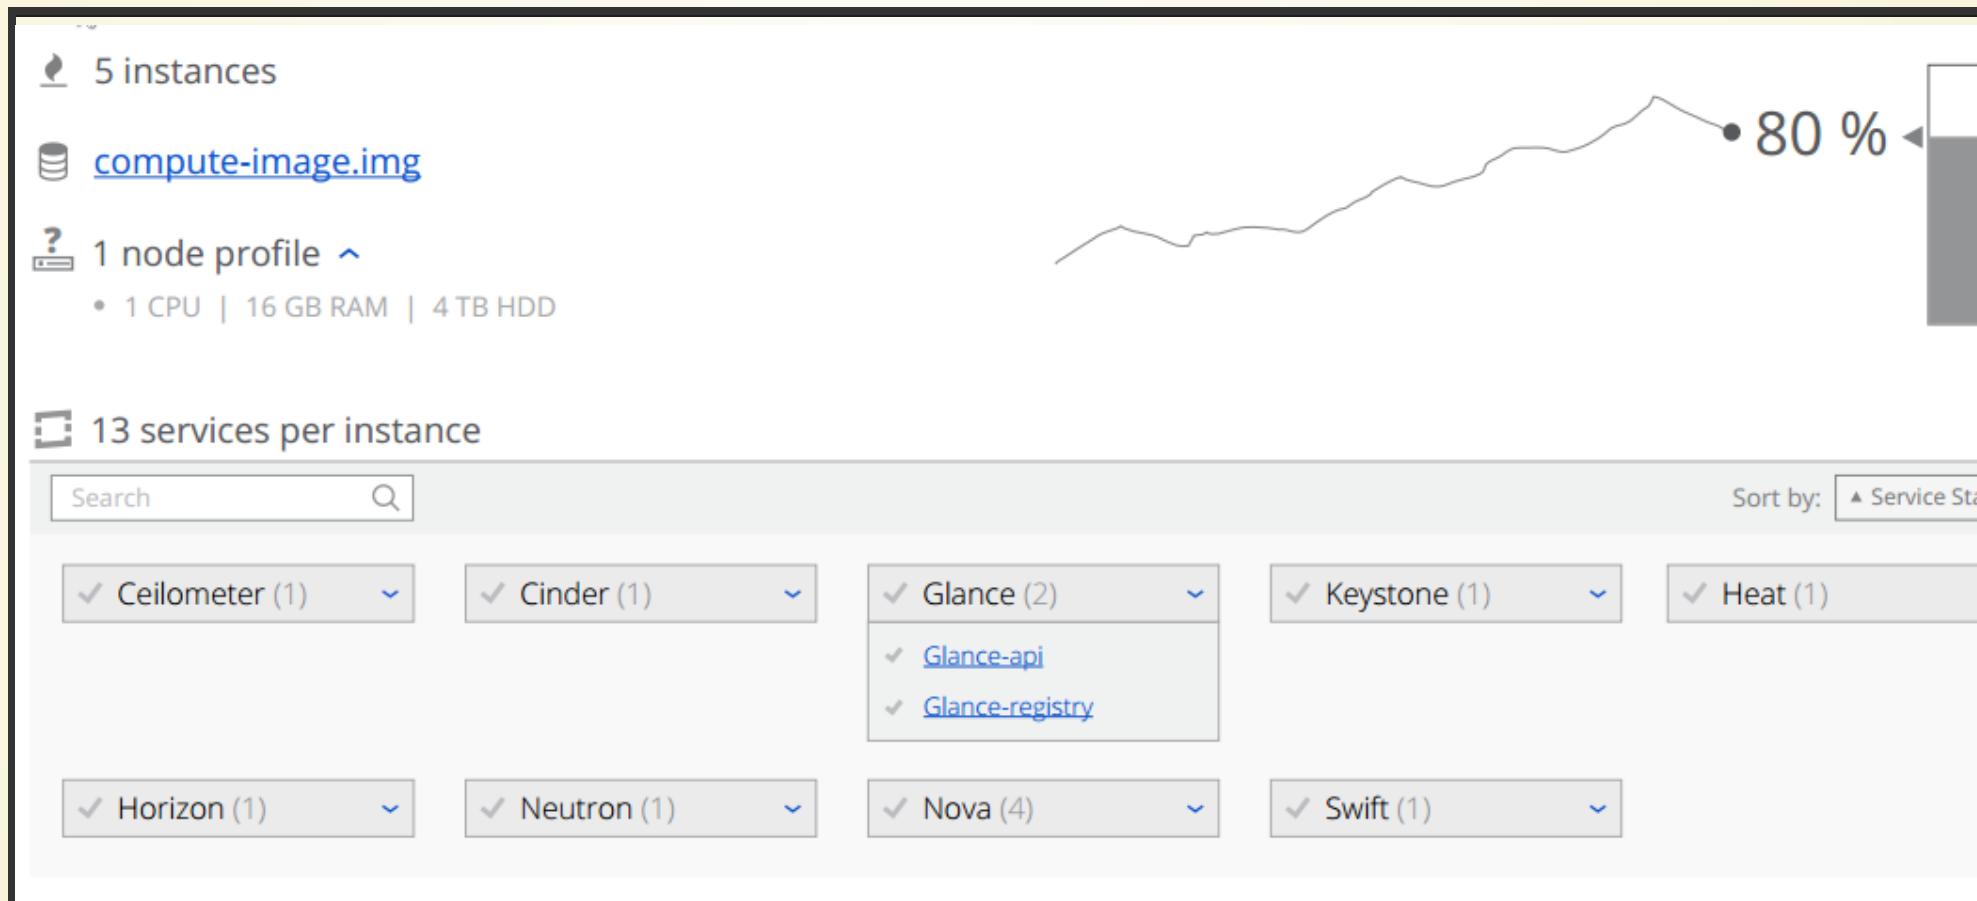
# FUTURE (JUNO) ADDING 'EVENTS TIME-LINE' TO CHARTS

- Under the each chart, there will be a time-line of the events, as a context for the chart data.
- E.g. machine updated, powered down, bad health, etc..

# FUTURE (JUNO) MONITORING OF RUNNING SERVICES

- Extending Glance, so it can return list of services packed in image
- Monitoring of services on each baremetal using Ceilometer

# FUTURE(JUNO) MONITORING OF THE RUNNING SERVICES



# OVERCLOUD MANAGEMENT WITH TUSKAR

# SCALING OF THE OVERCLOUD

## Scale Deployment

Deployment Role	Node Profile	Instances	Change
Controller	• Default	<div style="border: 1px solid #ccc; padding: 2px; text-align: center;">1</div>	no change
Compute	↑ My Node Profile • Different Profile	<div style="border: 1px solid #0070C0; padding: 2px; text-align: center;">55</div> <div style="display: flex; justify-content: space-around;"><span>◀</span><span>▶</span></div>	+5
Object Storage	• Default	<div style="border: 1px solid #ccc; padding: 2px; text-align: center;">10</div> <div style="display: flex; justify-content: space-around;"><span>◀</span><span>▶</span></div>	no change
Block Storage	• Default	<div style="border: 1px solid #ccc; padding: 2px; text-align: center;">4</div> <div style="display: flex; justify-content: space-around;"><span>◀</span><span>▶</span></div>	no change
		<div style="border: 1px solid #ccc; padding: 2px; text-align: center;">0</div> <div style="display: flex; justify-content: space-around;"><span>◀</span><span>▶</span></div>	no change

# DOWNSCALING OF THE OVERCLOUD

- Deleting resource manually then update stack.
- Choosing which resource to unprovision when downscaling, so users can evacuate services running inside, before downscaling happens
- Next: defining hooks for upscaling/downscaling so the evacuating behavior can be defined on template level, so it will be automatic.

# FUTURE (JUNO) AUTO-SCALING OF THE OVERCLOUD

- We are waiting on ability to scale nested stacks, implemented in Heat
- Auto-scaling using Ceilometer alarms defined in the Heat template
- Auto-scaling with conditions we need, e.g. `cpu_util` bigger than 50% for 1 minute

# FUTURE (JUNO) AUTO-SCALING EXAMPLES

- Newly plugged hardware can be auto-discovered and auto-scaled.
- It can be combined, as your application is scaling, the more hardware will be needed. That can be taken from e.g. other scaling down and releasing the hardware.

# RUNNING OVERCLOUD IN HA

- AMQP - Rabbit MQ cluster in active-active
- Openstack API instances - Virtual IP and keepalived + HA proxy
- DB - Mysql Percona XtraDB Cluster

# TUSKAR API AND TUSKAR-UI SUMMARY

- Deploy overcloud
- Monitor overcloud
- Manage overcloud
  - Scale
  - Update

# RESOURCES

- <https://wiki.openstack.org/wiki/TripleO>
- <http://docs.openstack.org/developer/tripleo-incubator/>

# QUESTIONS

**THANKS FOR YOUR ATTENTION!**  
**GIVE US FEEDBACK! [HTTP://DEVCONF.CZ/F/16](http://devconf.cz/f/16)**

